

2.4 Data Cube Operations

May 8, 2019

1 2.4 Data Cube Operations

Data cubes are d -dimensional hypercubes. We can answer questions about a data set by manipulating this hypercube. In this section, we will study three basic operations: slicing, dicing, and roll-ups.

```
In [1]: %matplotlib inline
import numpy as np
import pandas as pd
titanic_df = pd.read_csv("https://raw.githubusercontent.com/dlsun/data-science-book/master/titanic.csv")

titanic_df["adult"] = (titanic_df["age"] >= 18)

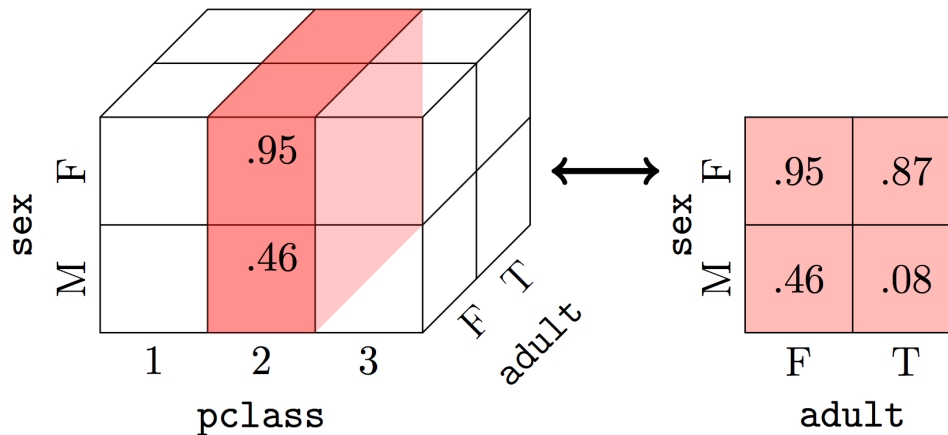
survival_cube = titanic_df.pivot_table(
    index="sex", columns=["pclass", "adult"],
    values="survived", aggfunc=np.mean)
survival_cube
```

```
Out[1]: pclass      1      2      3
adult    False    True    False    True    False    True
sex
female  0.947368  0.968000  0.952381  0.870588  0.536364  0.443396
male    0.400000  0.326389  0.464286  0.083916  0.147059  0.155709
```

2 Slicing

Slicing a data cube refers to fixing the value of one dimension of the hypercube. For example, suppose we only want to know the survival rates of passengers in second class. To do this, we fix the value of `pclass` at 2 and look at the survival rates over the other dimensions.

It is easy to see why this operation is called “slicing” if you imagine a three-dimensional cube. When we fix the value of one dimension, we are essentially slicing the cube at that value, as shown in the figure below.



Each slice reduces the dimension of a data cube by one. If the original data cube had d dimensions, then the slice has $d - 1$ dimensions.

To slice a pivot table in pandas, we simply access the corresponding row or column in the DataFrame. For example, to get the survival rates for the passengers in second class from the data cube above, we can simply select the column labeled 2. The result is a two-dimensional data cube:

```
In [2]: survival_cube[2]
```

```
Out[2]: adult      False      True
sex
female  0.952381  0.870588
male    0.464286  0.083916
```

Depending on how the pivot table is arranged, the slice is sometimes not in data cube form. For example, if we slice the data cube to get only the data for male passengers, the output is two-dimensional but not in data cube form:

```
In [3]: survival_cube.loc["male"]
```

```
Out[3]: pclass  adult
1      False    0.400000
        True     0.326389
2      False    0.464286
        True     0.083916
3      False    0.147059
        True     0.155709
Name: male, dtype: float64
```

But it is easy to convert this tabular data into a data cube; we simply unstack the Series so that each value of adult is a separate column.

```
In [4]: survival_cube.loc["male"].unstack()
```

```
Out[4]: adult      False      True
pclass
1      0.400000  0.326389
2      0.464286  0.083916
3      0.147059  0.155709
```

Slicing was easy in the two examples above because `pclass` and `sex` were both the outermost (i.e., first) level in their respective indexes. But what if we want to slice on a dimension that is buried in some intermediate level of a `MultiIndex`? We can use the “cross-section” function (`.xs`) of pandas. For example, the following code returns the survival rates for the children on the Titanic:

```
In [5]: survival_cube.xs(False, level="adult", axis=1)
```

```
Out[5]: pclass      1      2      3
       sex
       female  0.947368  0.952381  0.536364
       male    0.400000  0.464286  0.147059
```

This code tells pandas to return all columns (because `axis=1`) of `survival_cube` where `adult` is equal to `False`.

3 Dicing

Dicing is like slicing, except that we fix the values of two or more dimensions. For example, if we want to know the survival rates of males in second class, we could dice the data cube as follows:

```
In [6]: survival_cube.loc["male", 2]
```


```
Out[6]: adult
       False    0.464286
       True     0.083916
       Name: male, dtype: float64
```

Notice that the result is one-dimensional, since there is only one remaining dimension that we have not fixed (`adult`). In general, if we dice a d -dimensional data cube along k dimensions, the result will be a $(d - k)$ -dimensional data cube. In the example above, $d = 3$ and $k = 2$, so the output had $d - k = 1$ dimensions.

4 Roll-ups

The data cube above contains information about the survival rate by `pclass`, `sex`, and `adult`. But what if we are only interested in the survival rate by `sex` and `adult`? To do this, we have to **roll-up** data cube over the variable `pclass`.

The roll-up operation is diagrammed in the figure below. We want to collapse the `pclass` dimension, resulting in a two-dimensional data cube over `sex` and `adult`. Just as with the slicing operation, each roll-up operation reduces the dimension of the data cube by one.



What is the best way to calculate roll-ups? This is a trick question: the best way is to avoid calculating them at all! If we calculate the roll-ups when the pivot table is first created, then we can just look them up without having to calculate them. To get the roll-ups alongside the cell values, we specify `margins=True` in `.pivot_table()`.

```
In [7]: survival_cube_with_rollups = titanic_df.pivot_table(
        index=["sex", "adult"], columns="pclass",
        values="survived", aggfunc=np.mean,
        margins=True
    )
```

survival_cube_with_rollups

```
Out [7]: pclass      1      2      3      All
sex  adult
female False  0.947368  0.952381  0.536364  0.646667
      True   0.968000  0.870588  0.443396  0.765823
male   False  0.400000  0.464286  0.147059  0.213483
      True   0.326389  0.083916  0.155709  0.180556
All      0.619195  0.429603  0.255289  0.381971
```

Compared with the default pivot table (i.e., without `margins=True`), this pivot table has an extra row and an extra column, both labeled “All”. They are in the *margins* of the original table (which is why the extra argument to `.pivot_table()` was `margins=True`). This additional row and column contain various lower-dimensional roll-ups of the original data cube:

- The cell in the bottom right of this table is the roll-up of all three dimensions. In other words, it is the overall survival rate.
- The other values in the last column (labeled “All”) represent the roll-up of `pclass`. In other words, they are the survival rates by `sex` and `adult`.
- The other values in the last row (labeled “All”) represent the roll-up of both `sex` and `adult`. In other words, they are the survival rates by `pclass`.

However, this table does not store all the possible roll-ups of the three-dimensional datacube. For example, it does not store the roll-up of `sex`, nor does it store the roll-up of both `pclass` and `adult`. When designing a pivot table, it is a good idea to think about which roll-ups are most important and to choose the row indexes and columns accordingly so that those roll-ups are available.

4.1 Calculating Your Own Roll-Ups

Suppose you forgot to include the roll-ups when you first created the pivot table, or perhaps you need a roll-up that your pivot table does not provide. In most cases, there is no way to reconstruct the roll-ups from just the data cube. However, for some metrics, it is possible to reconstruct the roll-ups from the data cube.

For example, consider a data cube that stores the *number* of survivors by `sex`, `adult`, and `pclass`.

```
In [8]: num_survivors_cube = titanic_df.pivot_table(
        index=["sex", "adult"], columns="pclass",
        values="survived", aggfunc=np.sum
    )
```

```
num_survivors_cube
```

```
Out[8]: pclass      1    2    3
sex      adult
female False    18    20    59
         True   121    74    47
male    False    14    13    30
         True    47    12    45
```

Now, if we want to roll-up the `pclass` variable, we can calculate the total number of survivors by summing the numbers in first, second, and third class. In other words, we need to sum each row of the DataFrame above. This is possible using `.sum()`, but we have to specify an additional keyword argument, `axis=`, so that pandas knows which dimension to sum over:

- `axis=0` means aggregate *over* the rows (i.e., dimension 0), returning one number per column
- `axis=1` means aggregate *over* the columns (i.e., dimension 1), returning one number per row

Because we want the sum of each row, we are aggregating over the columns; thus we need to sum over `axis=1`:

```
In [9]: num_survivors_cube.sum(axis=1)
```

```
Out[9]: sex      adult
female False      97
         True   242
male    False      57
         True   104
dtype: int64
```

As a sanity check, let's make sure these numbers match the results from `.pivot_table()` when we set `margins=True`:

```
In [10]: titanic_df.pivot_table(
        index=["sex", "adult"], columns="pclass",
        values="survived", aggfunc=np.sum,
        margins=True
    )
```

```
Out[10]: pclass      1      2      3  All
sex      adult
female False   18    20    59   97
         True  121    74    47  242
male    False   14    13    30   57
         True   47    12    45  104
All                200  119  181  500
```

The numbers in the “All” column match exactly!

5 Exercises

Exercise 1. We saw one case where it was possible to manually reconstruct roll-ups using only the values in a data cube.

Is it possible to calculate the roll-up of `pclass` from just the values in `survival_cube` (a pivot table defined above)? In other words, can we reconstruct the survival rates by sex and adult from just the survival rates in `survival_cube`? Try a few different approaches and compare the results against the true answer, which you can obtain using `.groupby()` or `.pivot_table(..., margins=True)`.

```
In [23]: survival_cube
survival_cube.stack("adult").mean(axis=1), survival_cube.stack("adult").mean(axis=1).
#this method averages over p-class equally instead of proportionally across the number
```

```
Out[23]: (sex      adult
female  False    0.812038
         True     0.760661
male    False    0.337115
         True     0.188671
dtype: float64, adult      False      True
sex
female  0.812038  0.760661
male    0.337115  0.188671)
```

```
In [24]: titanic_df.groupby(["sex", "adult"])["survived"].mean()
#Using groupby will give us correct values; cannot calculate rollups after pivot table
```

```
Out[24]: sex      adult
female  False    0.646667
```

```

        True      0.765823
male    False      0.213483
        True      0.180556
Name: survived, dtype: float64

```

```

In [25]: titanic_df.pivot_table(
        index=["sex", "adult"], columns="pclass",
        values="survived", aggfunc="mean",
        margins=True)

```

#margins=TRUE gives us the correct margin averages by weight

```

Out[25]: pclass      1      2      3      All
sex    adult
female False  0.947368  0.952381  0.536364  0.646667
        True   0.968000  0.870588  0.443396  0.765823
male   False  0.400000  0.464286  0.147059  0.213483
        True   0.326389  0.083916  0.155709  0.180556
All                0.619195  0.429603  0.255289  0.381971

```

```

In [26]: titanic_df.pivot_table(
        index=["sex", "adult"], columns="pclass",
        values="survived", aggfunc="count",
        margins=True)

```

```

Out[26]: pclass      1      2      3      All
sex    adult
female False    19    21   110   150
        True   125    85   106   316
male   False    35    28   204   267
        True   144   143   289   576
All                323   277   709  1309

```

Exercises 2-4 deal with the Tips data set (<https://raw.githubusercontent.com/dlsun/data-science-book/master/tips.csv>)

Exercise 2. Create a pivot table that shows the average total bill by day, time, and table size. Include roll-ups with this pivot table that make it easy to answer questions like, “Is the average bill higher for lunch or dinner?”

```

In [12]: # TYPE YOUR CODE HERE.

```

```

tips = pd.read_csv("https://raw.githubusercontent.com/dlsun/data-science-book/master/tips.csv")

avg_total_bill_pivot = tips.pivot_table(
    index=["time"], columns=["size", "day"],
    values="total_bill", aggfunc=np.mean,
    margins=True)

avg_total_bill_pivot

```

```

Out[12]: size      1      2      3 \
day      Fri  Sat  Thur  Fri  Sat  Sun  Thur  Fri

```

time									
Dinner	NaN	5.16	NaN	17.799091	16.83717	17.56	18.780000	NaN	
Lunch	8.58	NaN	10.07	13.072000	NaN	NaN	15.079787	15.98	
All	8.58	5.16	10.07	16.321875	16.83717	17.56	15.156875	15.98	

size				...	4				\
day		Sat	Sun	...	Fri	Sat	Sun	Thur	
time				...					
Dinner	25.509444	22.184	...	40.17	29.876154	26.688333	NaN		
Lunch	NaN	NaN	...	NaN	NaN	NaN	29.95		
All	25.509444	22.184	...	40.17	29.876154	26.688333	29.95		

size	5			6		All			
day	Sat	Sun	Thur	Sun	Thur				
time									
Dinner	28.15	27.0	NaN	48.17	NaN	20.797159			
Lunch	NaN	NaN	41.19	NaN	30.383333	17.168676			
All	28.15	27.0	41.19	48.17	30.383333	19.785943			

[3 rows x 21 columns]

Exercise 3. Create a pivot table that shows that average total bill by day and time for parties of size 2. (Don't do this by calling `.pivot_table()` on the original data. You should be able to do this using just the pivot table you created in Exercise 2.)

```
In [13]: # TYPE YOUR CODE HERE.
avg_total_bill_pivot[2]

#if size and day was flipped, use avg_total_bill_pivot.xs(2, level="size", axis=1)
```

```
Out[13]: day          Fri          Sat          Sun          Thur
time
Dinner  17.799091  16.83717  17.56  18.780000
Lunch   13.072000          NaN          NaN  15.079787
All     16.321875  16.83717  17.56  15.156875
```

Exercise 4. How would you create a pivot table that shows the average total bill by day and time? Is it possible to do this using just the pivot table you created in Exercise 2?

```
In [14]: avg_total_bill_pivot
```

```
Out[14]: size          1          2          3 \
day          Fri  Sat  Thur          Fri  Sat  Sun          Thur  Fri
time
Dinner  NaN  5.16  NaN  17.799091  16.83717  17.56  18.780000  NaN
Lunch   8.58  NaN  10.07  13.072000          NaN  NaN  15.079787  15.98
All     8.58  5.16  10.07  16.321875  16.83717  17.56  15.156875  15.98

size          ...          4 \
```


day	Sat	Sun	...	Fri	Sat	Sun	Thur
time			...				
Dinner	25.509444	22.184	...	40.17	29.876154	26.688333	NaN
Lunch	NaN	NaN	...	NaN	NaN	NaN	29.95
All	25.509444	22.184	...	40.17	29.876154	26.688333	29.95

size	5		6		All
day	Sat	Sun	Thur	Sun	Thur
time					
Dinner	28.15	27.0	NaN	48.17	NaN
Lunch	NaN	NaN	41.19	NaN	30.383333
All	28.15	27.0	41.19	48.17	30.383333

[3 rows x 21 columns]

```
In [15]: # TYPE YOUR CODE HERE.
```

```
(avg_total_bill_pivot.drop("All", axis=0).drop("All", axis=1)).stack("day").mean(axis=1)
```

```
Out[15]: time    day
Dinner  Fri      28.984545
        Sat      21.106554
        Sun      28.320467
        Thur     18.780000
Lunch   Fri      12.544000
        Thur     24.305520
dtype: float64
```

```
In [16]: avg_total_bill_pivot.stack("day").mean(axis=1)
```

```
Out[16]: time    day
Dinner          20.797159
        Fri      28.984545
        Sat      21.106554
        Sun      28.320467
        Thur     18.780000
Lunch          17.168676
        Fri      12.544000
        Thur     24.305520
All            19.785943
        Fri      20.262969
        Sat      21.106554
        Sun      28.320467
        Thur     24.318368
dtype: float64
```

```
In [17]: tips.groupby(["time", "day"]).total_bill.mean()
```

*#This method is correct because in avg_total_bill_pivot (data cube), the method we use
#the average of an average which is incorrect. Each entry in the data cube is average
#given equal weight which is not valid.*

```
Out[17]: time    day
         Dinner  Fri    19.663333
          Sat    20.441379
          Sun    21.410000
          Thur    18.780000
         Lunch  Fri    12.845714
          Thur    17.664754
         Name: total_bill, dtype: float64
```