

5.5 Model Selection and Hyperparameter Tuning

May 9, 2019

1 5.5 Model Selection and Hyperparameter Tuning

This section will use the tools developed in the previous section to answer two important questions:

- Model Selection: How do we determine which model is best?
- Hyperparameter Tuning: How do we choose hyperparameters, such as k in k -nearest neighbors?

In the previous section, we saw how to use training and validation sets to estimate how well the model will perform on future data. A natural way to decide between competing models (or hyperparameters) is to choose the one that minimizes the validation error.

```
In [1]: %matplotlib inline
import numpy as np
import pandas as pd
pd.options.display.max_rows = 5

housing = pd.read_csv("https://raw.githubusercontent.com/dlsun/data-science-book/master/
                    sep="\t")

housing
```

```
Out[1]:
```

	Order	PID	MS SubClass	MS Zoning	Lot Frontage	Lot Area	Street	\
0	1	526301100	20	RL	141.0	31770	Pave	
1	2	526350040	20	RH	80.0	11622	Pave	
...	
2928	2929	924100070	20	RL	77.0	10010	Pave	
2929	2930	924151050	60	RL	74.0	9627	Pave	

	Alley	Lot Shape	Land Contour	...	Pool Area	Pool QC	Fence	\
0	NaN	IR1	Lvl	...	0	NaN	NaN	
1	NaN	Reg	Lvl	...	0	NaN	MnPrv	
...	
2928	NaN	Reg	Lvl	...	0	NaN	NaN	
2929	NaN	Reg	Lvl	...	0	NaN	NaN	

Misc Feature	Misc Val	Mo Sold	Yr Sold	Sale Type	Sale Condition	\
--------------	----------	---------	---------	-----------	----------------	---

0	NaN	0	5	2010	WD	Normal
1	NaN	0	6	2010	WD	Normal
...
2928	NaN	0	4	2006	WD	Normal
2929	NaN	0	11	2006	WD	Normal


	SalePrice
0	215000
1	105000
...	...
2928	170000
2929	188000

[2930 rows x 82 columns]

1.1 K-Fold Cross Validation

Previously, we carried out cross validation by splitting the data into 2 halves, alternately using one half to train the model and the other to evaluate the model. In general, we can split the data into k subsamples, alternately training the data on $k - 1$ subsamples and evaluating the model on the 1 remaining subsample, i.e., the validation set. This produces k somewhat independent estimates of the test error. This procedure is known as **k -fold cross validation**. (Be careful not to confuse the k in k -fold cross validation with the k in k -nearest neighbors.) Therefore, the specific version of cross validation that we saw earlier is 2-fold cross validation.

A schematic of 4-fold cross validation is shown below.



Implementing k -fold cross validation from scratch for $k > 2$ is straightforward but messy, so we will usually let Scikit-Learn do it for us.

1.2 Cross Validation in Scikit-Learn

Scikit-Learn provides a function, `cross_val_score`, that will carry out all aspects of k -fold cross validation:

1. split the data into k subsamples
2. combine the first $k - 1$ subsamples into a training set and train the model
3. evaluate the model predictions on the last (k th) held-out subsample
4. repeat steps 2-3 k times (i.e. k “folds”), each time holding out a different one of the k subsamples

5. calculate k “scores”, one from each validation set

There is one subtlety to keep in mind. Training a k -nearest neighbors model is not just about fitting the model; it also involves dummifying the categorical variables and scaling the variables. These preprocessing steps should be included in the cross-validation process. They cannot be done ahead of time.

For example, suppose we run 5-fold cross validation. Then:

- When subsamples 1-4 are used for training and subsample 5 for validation, the observations have to be standardized with respect to the mean and SD of subsamples 1-4.
- When subsamples 2-5 are used for training and subsample 1 for validation, the observations have to be standardized with respect to the mean and SD of subsamples 2-5.
- And so on.

We cannot simply standardize all of the data once at the beginning and run cross validation on the standardized data. To do so would be allowing the model to peek at the validation set during training. That’s because each training set would be standardized with respect to a mean and SD that is calculated from all data, including the validation set. To be completely above board, we should standardize each training set with respect to the mean and SD of just that training set.

Fortunately, Scikit-Learn provides a Pipeline object that allows us to chain these preprocessing steps together with the model we want to fit.

```
In [2]: from sklearn.feature_extraction import DictVectorizer
        from sklearn.preprocessing import StandardScaler
        from sklearn.neighbors import KNeighborsRegressor
        from sklearn.pipeline import Pipeline

        # get the features (in dict format) and the labels
        # (do not split into training and validation sets)
        features = ["Lot Area", "Gr Liv Area",
                    "Full Bath", "Half Bath",
                    "Bedroom AbvGr",
                    "Year Built", "Yr Sold",
                    "Neighborhood"]
        X_dict = housing[features].to_dict(orient="records")
        y = housing["SalePrice"]

        # specify the pipeline
        vec = DictVectorizer(sparse=False)
        scaler = StandardScaler()
        model = KNeighborsRegressor(n_neighbors=10)
        pipeline = Pipeline([("vectorizer", vec), ("scaler", scaler), ("fit", model)])
```

This entire Pipeline can be passed to `cross_val_score`, along with the data, the number of folds k (cv), and the type of score (scoring). So 5-fold cross validation in Scikit-Learn would look as follows:

```
In [3]: from sklearn.model_selection import cross_val_score
```

```
scores = cross_val_score(pipeline, X_dict, y,
                          cv=5, scoring="neg_mean_squared_error")
scores
```

```
Out [3]: array([-1.98025092e+09, -1.47950953e+09, -1.54632969e+09,
               -1.87662563e+09, -1.42167566e+09])
```

Notice that we get five (negative) validation MSEs, one from each of the 5 folds. `cross_val_score` returns the *negative* MSE, instead of the MSE, because by definition, a *higher* score is better. (Since we want the MSE to be as *low* as possible, we want the negative MSE to be as *high* as possible.)

To come up with a single overall estimate of the test MSE, we flip the signs and average the MSEs:

```
In [4]: np.mean(-scores)
```

```
Out [4]: 1660878287.4356349
```

The RMSE is the square root of the MSE:

```
In [5]: np.sqrt(np.mean(-scores))
```

```
Out [5]: 40753.874508267734
```

1.3 Hyperparameter Tuning

How do we choose k ? We can simply try all values of k and pick the one with the smallest (test) MSE.


```
In [6]: vec = DictVectorizer(sparse=False)
        scaler = StandardScaler()

        # calculates estimate of test error based on 10-fold cross validation
        def get_cv_error(k):
            model = KNeighborsRegressor(n_neighbors=k)
            pipeline = Pipeline([("vectorizer", vec), ("scaler", scaler), ("fit", model)])
            mse = np.mean(-cross_val_score(
                pipeline, X_dict, y,
                cv=10, scoring="neg_mean_squared_error"
            ))
            return mse

        ks = pd.Series(range(1, 51))
        ks.index = range(1, 51)
        test_errs = ks.apply(get_cv_error)

        test_errs.plot.line()
        test_errs.sort_values()
```

```
Out [6]: 4      1.390511e+09
        3      1.405909e+09
        ...
        49     2.208727e+09
        50     2.216495e+09
Length: 50, dtype: float64
```



The MSE is minimized near $k = 4$, which suggests that a 4-nearest neighbors model is optimal for prediction.

1.4 Model Selection

Suppose we are not sure whether Yr Sold should be included in the 4-nearest neighbors model or not. To determine whether or not it should be included, we can fit a model with Yr Sold included and another model with it excluded, and see which model has the better (test) MSE.

```
In [7]: vec = DictVectorizer(sparse=False)
        scaler = StandardScaler()
        model = KNeighborsRegressor(n_neighbors=4)
        pipeline = Pipeline([("vectorizer", vec), ("scaler", scaler), ("fit", model)])
```

```
In [8]: features = ["Lot Area", "Gr Liv Area",
                   "Full Bath", "Half Bath",
                   "Bedroom AbvGr",
                   "Year Built", "Yr Sold",
                   "Neighborhood"]
```

```
X_dict = housing[features].to_dict(orient="records")
np.mean(
    -cross_val_score(pipeline, X_dict, y, cv=10, scoring="neg_mean_squared_error")
)
```

Out[8]: 1390511354.3758318

```
In [9]: features = ["Lot Area", "Gr Liv Area",
                    "Full Bath", "Half Bath",
                    "Bedroom AbvGr",
                    "Year Built",
                    "Neighborhood"]
X_dict = housing[features].to_dict(orient="records")
-cross_val_score(pipeline, X_dict, y, cv=10, scoring="neg_mean_squared_error").mean()
```

Out[9]: 1315348887.0818686

The MSE actually goes down when we remove Yr Sold, so it seems that the model is better off without this variable.