

11.2 The XML Data Format

May 9, 2019

1 11.2 The XML Data Format

XML, which stands for eXtensible Markup Language, is another way to represent hierarchical data. The basic building block of XML is the **tag**, denoted by angle brackets <>.

For example, a data set of movies might be represented using XML as follows:

```
<movies>
  <movie id="1" title="The Godfather">
    <director id="50" name="Coppola, Francis Ford">
    </director>
    <releasedate>1972-03-24</releasedate>
    <character id="100" name="Vito Corleone">
      <actor id="200" name="Brando, Marlon">
      </actor>
    </character>
    <character id="101" name="Michael Corleone">
      <actor id="201" name="Pacino, Al">
      </actor>
    </character>
    ...
  </movie>
  <movie id="2" title="The Godfather: Part II">
    <director id="50" name="Coppola, Francis Ford">
    </director>
    <releasedate>1974-10-20</releasedate>
    <character id="101" name="Michael Corleone">
      <actor id="201" name="Pacino, Al">
      </actor>
    </character>
    <character id="100" name="Vito Corleone">
      <actor id="250" name="De Niro, Robert">
      </actor>
    </character>
    ...
  </movie>
  ...
</movies>
```

Note the following features of XML:

- Every tag `<a>` has a corresponding closing tag ``. You can always recognize a closing tag by the forward slash `/`.
- Additional tags and/or strings can be nested between the opening and closing tags. In the example above, `<actor>` is nested between `<character>` and `</character>`, and `<character>` is nested between `<movie>` and `</movie>`. The nesting is used to represent hierarchy.
- Indentation is used to make the code more readable (to make it easier to see the nesting structure). But it is optional.
- Attributes can be associated with each tag, like `id=` and `name=` with the `<character>` tag and `id=` and `title=` with the `<movie>` tag.

Each tag represents a variable in the data set. Unlike JSON, which uses lists to represent repeated fields, XML represents repeated fields by simply repeating tags where necessary. In the example above, there are multiple instances of `<movie>` within `<movies>` and multiple instances of `<character>` within `<movie>`, so `movie` and `character` are both repeated fields. (In fact, `director` is also a repeated field, but it is impossible to tell from the code above, since the movies shown above only have one director.)

You will learn XML by working with the same New York Philharmonic data as in the previous section, except that the data is now stored in XML format. Let's look at this file on disk.

```
In [1]: !ls -l /data301/data/nyphil/
```

```
total 78912
-rw-r--r-- 1 root root 35284693 Feb 25 08:35 complete.json
-rw-r--r-- 1 root root 45514892 Feb 25 08:36 complete.xml
```

Notice that this XML file is nearly twice as large as the JSON file. Although XML is more readable than JSON, it is a more expensive way to store hierarchical data, primarily because of the cost of storing both the opening and closing tags.

There are several libraries in Python for working with XML, including BeautifulSoup (which we will use in the next section to parse HTML), ElementTree, and lxml. We will use lxml to work with XML data because it is fastest for large data sets, provided that the data is well-formed. The lxml library provides a convenient API that replicates all of the functionality of ElementTree, plus implements a few additional features that are useful for data analysis.

```
In [2]: from lxml import etree
```

First, let's read in the data using lxml. The `.parse()` function of ElementTree reads in an XML document from a file or URL and returns a tree-like data structure called an ElementTree.

```
In [3]: tree = etree.parse("/data301/data/nyphil/complete.xml")
```

Every XML document has a single "root" tag that encloses all of the other tags. For the New York Philharmonic data, this root tag is `<programs>`.

```
In [4]: tree.getroot()
```

```
Out[4]: <Element programs at 0x7f4a50d5dc48>
```

If the XML data is already stored as a string in memory, then we instead use the `.fromstring()` method. Note that `.fromstring()` returns the root tag directly.

```
In [5]: with open("/data301/data/nyphil/complete.xml", "rb") as f:
        string = f.read()

        etree.fromstring(string)
```

```
Out[5]: <Element programs at 0x7f4a50060ac8>
```

Each direct descendant, or **child**, of `<programs>` is a program. To find the direct descendants of a tag, we call the `.getchildren()` method.

```
In [6]: programs = tree.getroot()
        print(len(programs.getchildren()))
        programs.getchildren()[:10]
```

```
14009
```

```
Out[6]: [<Element program at 0x7f4a39c0c688>,
        <Element program at 0x7f4a39c0c708>,
        <Element program at 0x7f4a39c0c108>,
        <Element program at 0x7f4a39c0c908>,
        <Element program at 0x7f4a39c0c948>,
        <Element program at 0x7f4a39c0ca08>,
        <Element program at 0x7f4a39c0ca48>,
        <Element program at 0x7f4a39c0ca88>,
        <Element program at 0x7f4a39c0c988>,
        <Element program at 0x7f4a39c0c9c8>]
```

Let's print out the first of these programs. There are two ways to get the first program.

```
In [7]: # METHOD 1: Get it from the list above.
        program = programs.getchildren()[0]

        # METHOD 2: Use .find() to find the first instance of a tag.
        program = tree.find("program")
        program
```

```
Out[7]: <Element program at 0x7f4a39c0c688>
```

Now let's see how the data is represented by printing out the XML of this program. To do this, we use the `etree.tostring()` function.

```
In [8]: print(etree.tostring(program, encoding="unicode"))
```

```

<program>
  <id>00646b9f-fec7-4ffb-9fb1-faae410bd9dc-0.1</id>
  <programID>3853</programID>
  <orchestra>New York Philharmonic</orchestra>
  <season>1842-43</season>
  <concertInfo>
    <eventType>Subscription Season</eventType>
    <Location>Manhattan, NY</Location>
    <Venue>Apollo Rooms</Venue>
    <Date>1842-12-07T05:00:00Z</Date>
    <Time>8:00PM</Time>
  </concertInfo>
  <worksInfo>
    <work ID="52446*">
      <composerName>Beethoven, Ludwig van</composerName>
      <workTitle>SYMPHONY NO. 5 IN C MINOR, OP.67</workTitle>
      <conductorName>Hill, Ureli Corelli</conductorName>
    </work>
    <work ID="8834*4">
      <composerName>Weber, Carl Maria Von</composerName>
      <workTitle>OBERON</workTitle>
      <movement>"Ozean, du Ungeheuer" (Ocean, thou mighty monster), Reiza (Scene and
      <conductorName>Timm, Henry C.</conductorName>
      <soloists>
        <soloist>
          <soloistName>Otto, Antoinette</soloistName>
          <soloistInstrument>Soprano</soloistInstrument>
          <soloistRoles>S</soloistRoles>
        </soloist>
      </soloists>
    </work>
    <work ID="3642*">
      <composerName>Hummel, Johann</composerName>
      <workTitle>QUINTET, PIANO, D MINOR, OP. 74</workTitle>
      <soloists>
        <soloist>
          <soloistName>Scharfenberg, William</soloistName>
          <soloistInstrument>Piano</soloistInstrument>
          <soloistRoles>A</soloistRoles>
        </soloist>
        <soloist>
          <soloistName>Hill, Ureli Corelli</soloistName>
          <soloistInstrument>Violin</soloistInstrument>
          <soloistRoles>A</soloistRoles>
        </soloist>
        <soloist>
          <soloistName>Derwort, G. H.</soloistName>
          <soloistInstrument>Viola</soloistInstrument>

```

```

        <soloistRoles>A</soloistRoles>
    </soloist>
    <soloist>
        <soloistName>Boucher, Alfred</soloistName>
        <soloistInstrument>Cello</soloistInstrument>
        <soloistRoles>A</soloistRoles>
    </soloist>
    <soloist>
        <soloistName>Rosier, F. W.</soloistName>
        <soloistInstrument>Double Bass</soloistInstrument>
        <soloistRoles>A</soloistRoles>
    </soloist>
</soloists>
</work>
<work ID="0*">
    <interval>Intermission</interval>
</work>
<work ID="8834*3">
    <composerName>Weber, Carl Maria Von</composerName>
    <workTitle>OBERON</workTitle>
    <movement>Overture</movement>
    <conductorName>Etienne, Denis G.</conductorName>
</work>
<work ID="8835*1">
    <composerName>Rossini, Gioachino</composerName>
    <workTitle>ARMIDA</workTitle>
    <movement>Duet</movement>
    <conductorName>Timm, Henry C.</conductorName>
    <soloists>
        <soloist>
            <soloistName>Otto, Antoinette</soloistName>
            <soloistInstrument>Soprano</soloistInstrument>
            <soloistRoles>S</soloistRoles>
        </soloist>
        <soloist>
            <soloistName>Horn, Charles Edward</soloistName>
            <soloistInstrument>Tenor</soloistInstrument>
            <soloistRoles>S</soloistRoles>
        </soloist>
    </soloists>
</work>
<work ID="8837*6">
    <composerName>Beethoven, Ludwig van</composerName>
    <workTitle>FIDELIO, OP. 72</workTitle>
    <movement>"In Des Lebens Fruhlingstagen...0 spur ich nicht linde," Florestan (
    <conductorName>Timm, Henry C.</conductorName>
    <soloists>
        <soloist>

```

```

        <soloistName>Horn, Charles Edward</soloistName>
        <soloistInstrument>Tenor</soloistInstrument>
        <soloistRoles>S</soloistRoles>
    </soloist>
</soloists>
</work>
<work ID="8336*4">
    <composerName>Mozart, Wolfgang Amadeus</composerName>
    <workTitle>ABDUCTION FROM THE SERAGLIO, THE, K.384</workTitle>
    <movement>"Ach Ich liebte," Konstanze (aria)</movement>
    <conductorName>Timm, Henry C.</conductorName>
    <soloists>
        <soloist>
            <soloistName>Otto, Antoinette</soloistName>
            <soloistInstrument>Soprano</soloistInstrument>
            <soloistRoles>S</soloistRoles>
        </soloist>
    </soloists>
</work>
<work ID="5543*">
    <composerName>Kalliwoda, Johann W.</composerName>
    <workTitle>OVERTURE NO. 1, D MINOR, OP. 38</workTitle>
    <conductorName>Timm, Henry C.</conductorName>
</work>
</worksInfo>
</program>

```

Hopefully, the basic structure of this data is already familiar to you from previous section. “Work”, “concertInfo”, and “soloist” are repeated fields inside “program”. One difference between the JSON and the XML is that “work” is not directly nested within “program”; the “work” tags are all nested inside an additional “worksInfo” tag.

Now suppose that we want to flatten the data at the level of soloists. To get all of the soloists, we can use the `.findall()` method. Let’s first try the obvious solution, which does not work:

```
In [9]: programs.findall("soloist")
```

```
Out[9]: []
```

Why did `lxml` fail to find any `<soloist>` tags? That’s because `.findall()` only searches among the direct descendants of a tag. We called `.findall()` on the `<programs>` tag, but all of its descendants are `<program>` tags.

To specify that `lxml` should look for `<soloist>` tags among all descendants, not just direct ones, we use the `.xpath()` command, which allows us to specify an XPath expression. [XPath](#) is a language used to select nodes from XML documents. The XPath expression to select all descendants named `<soloist>` of the current tag is `"//soloist"`. We pass this expression to the `.xpath()` method.

```
In [10]: soloists = programs.xpath("./soloist")
        len(soloists)
```

```
Out[10]: 56931
```

Now, to flatten the data at the level of soloists, we just need to turn `soloists` into a `DataFrame` with as many rows. But what if we want to include information from parent levels, like the composer of the work the soloist played? There are two ways.

1.0.1 Method 1

Since `<composerName>` is a descendant of `<work>`, one way is to navigate up to the level of `<work>` by calling `.getparent()` repeatedly and then find `<composerName>` among its descendants:

```
In [11]: soloist = soloists[0]

        # The first .getparent() returns the <soloists> tag.
        # The second .getparent() returns the <work> tag.
        # You have to figure this out by inspecting the XML.
        work = soloist.getparent().getparent()
        work.xpath("./composerName")
```

```
Out[11]: [<Element composerName at 0x7f4a50d5de88>]
```

This is a list with one tag, so we extract that tag and the text inside it.

```
In [12]: work.xpath("./composerName")[0].text
```

```
Out[12]: 'Weber, Carl Maria Von'
```

1.0.2 Method 2

As the number of levels of nesting increases, it quickly becomes impractical to call `.getparent()` repeatedly. We want to be able to jump directly to the right ancestor. The easiest way to do this is to use the XPath expression for an ancestor. To search for all ancestors named “work”, we can use the XPath expression `"ancestor::work"`.

```
In [13]: soloist.xpath("ancestor::work")
```

```
Out[13]: [<Element work at 0x7f4a398a0a88>]
```

Now, we can extract this single work tag and find its descendants named `<composerName>`. Or better yet, we can combine this step with the above step into a single XPath expression.

```
In [14]: soloist.xpath("ancestor::work//composerName")[0].text
```

```
Out[14]: 'Weber, Carl Maria Von'
```

Now let’s put it all together. We will flatten the data to get a `DataFrame` with one soloist per row. We will keep track of the soloist’s name, instrument, and role—as well as the composer of the work they performed. Unfortunately, it is much more manual to do this with XML than with JSON. There is no XML equivalent of the `json_normalize` function that will automatically produce a `DataFrame`, so we have to construct the `DataFrame` ourselves.

```
In [15]: import pandas as pd
```

```
rows = []
```

```
soloists = programs.xpath("./soloist")
```

```
for soloist in soloists:
```

```
    row = {}
```

```
    row["soloistName"] = soloist.find("soloistName").text
```

```
    row["soloistInstrument"] = soloist.find("soloistInstrument").text
```

```
    row["soloistRoles"] = soloist.find("soloistRoles").text
```

```
    row["composerName"] = soloist.xpath("ancestor::work//composerName")[0].text
```

```
    rows.append(row)
```

```
soloistsdf = pd.DataFrame(rows)
```

```
soloistsdf
```

```
Out[15]:
```

	composerName	soloistInstrument	\
0	Weber, Carl Maria Von	Soprano	
1	Hummel, Johann	Piano	
2	Hummel, Johann	Violin	
3	Hummel, Johann	Viola	
4	Hummel, Johann	Cello	
5	Hummel, Johann	Double Bass	
6	Rossini, Gioachino	Soprano	
7	Rossini, Gioachino	Tenor	
8	Beethoven, Ludwig van	Tenor	
9	Mozart, Wolfgang Amadeus	Soprano	
10	Bellini, Vincenzo	Soprano	
11	Romberg, Bernhard	Cello	
12	Rossini, Gioachino	Soprano	
13	Hummel, Johann	Piano	
14	Hummel, Johann	Piano	
15	Weber, Carl Maria Von	Soprano	
16	Weber, Carl Maria Von	Piano	
17	Hummel, Johann	Piano	
18	Pacini, Giovanni	Soprano	
19	Pacini, Giovanni	Piano	
20	Romberg, Bernhard	Cello	
21	Onslow, George	Piano	
22	Onslow, George	Flute	
23	Onslow, George	Clarinet	
24	Onslow, George	Bassoon	
25	Onslow, George	French Horn	
26	Onslow, George	Double Bass	
27	Onslow, George	None	
28	Onslow, George	Piano	
29	Onslow, George	Flute	
...	

56901	Klein, Gideon	Violin
56902	Klein, Gideon	Viola
56903	Klein, Gideon	Cello
56904	Beethoven, Ludwig van	Oboe
56905	Beethoven, Ludwig van	Oboe
56906	Beethoven, Ludwig van	Clarinet
56907	Beethoven, Ludwig van	Clarinet
56908	Beethoven, Ludwig van	Bassoon
56909	Beethoven, Ludwig van	Bassoon
56910	Beethoven, Ludwig van	French Horn
56911	Beethoven, Ludwig van	French Horn
56912	Shostakovich, Dmitri	Violin
56913	Shostakovich, Dmitri	Violin
56914	Shostakovich, Dmitri	Viola
56915	Shostakovich, Dmitri	Cello
56916	Shostakovich, Dmitri	Piano
56917	Mozart, Wolfgang Amadeus	Oboe
56918	Mozart, Wolfgang Amadeus	Clarinet
56919	Mozart, Wolfgang Amadeus	Bassoon
56920	Mozart, Wolfgang Amadeus	French Horn
56921	Handel, George Frideric	Soprano
56922	Handel, George Frideric	Mezzo-Soprano
56923	Handel, George Frideric	Tenor
56924	Handel, George Frideric	Bass-Baritone
56925	Handel, George Frideric	Chorus
56926	Handel, George Frideric	Soprano
56927	Handel, George Frideric	Mezzo-Soprano
56928	Handel, George Frideric	Tenor
56929	Handel, George Frideric	Baritone
56930	Handel, George Frideric	Chorus

	soloistName	soloistRoles
0	Otto, Antoinette	S
1	Scharfenberg, William	A
2	Hill, Ureli Corelli	A
3	Derwort, G. H.	A
4	Boucher, Alfred	A
5	Rosier, F. W.	A
6	Otto, Antoinette	S
7	Horn, Charles Edward	S
8	Horn, Charles Edward	S
9	Otto, Antoinette	S
10	Otto, Antoinette	S
11	Boucher, Alfred	S
12	Otto, Antoinette	S
13	Timm, Henry C.	S
14	Timm, Henry C.	S
15	Otto, Antoinette	S

16	Timm, Henry C.	A
17	Scharfenberg, William	S
18	Otto, Antoinette	S
19	Timm, Henry C.	A
20	Boucher, Alfred	S
21	Scharfenberg, William	S
22	Lehman	A
23	Groneveldt, Theodore W.	A
24	Hegelund, H. W.	A
25	Woehning, F. C.	A
26	Rosier, F. W.	A
27	None	None
28	Scharfenberg, William	S
29	Lehman	A
...
56901	Ge, Quan	A
56902	Young, Rebecca	A
56903	Gonzales, Alexei Yupanqui	A
56904	Sylar, Sherry	A
56905	Botti, Robert	A
56906	Martinez [Martínez] Forteza, Pascual	A
56907	Zoloto, Amy	A
56908	Laskowski, Kim	A
56909	Fast, Arlen	A
56910	Deane, Richard	A
56911	Spanjer, R. Allen	A
56912	Yao, Shanshan	A
56913	Rossano, Marié	A
56914	Kenote, Peter	A
56915	Tu, Qiang	A
56916	Wolfram, William	A
56917	Wang, Liang	S
56918	McGill, Anthony	S
56919	LeClair, Judith	S
56920	Deane, Richard	S
56921	Harvey, Joelle [Joéllle]	S
56922	Johnson Cano, Jennifer	S
56923	Bliss, Ben	S
56924	Foster-Williams, Andrew	S
56925	Westminster Symphonic Choir	S
56926	Harvey, Joelle [Joéllle]	S
56927	Johnson Cano, Jennifer	S
56928	Bliss, Ben	S
56929	Duncan, Tyler	S
56930	Westminster Symphonic Choir	S

[56931 rows x 4 columns]

Now, this is a DataFrame that we can analyze easily. For example, here is how many times Benny Goodman programmed a work by Mozart with the NY Phil:

```
In [16]: soloistsdf[soloistsdf["soloistName"] == "Goodman, Benny"].composerName.value_counts()
```

```
Out[16]: Mozart, Wolfgang Amadeus      3
         Weber, Carl Maria Von         3
         Gershwin, George              2
         Sauter, Eddie                 2
         Basie, Count                  1
         Williams, Mary Lou            1
         Prima, Louis                  1
         Youmans, Vincent              1
         Copland, Aaron                1
         Unspecified,                  1
         Green, Johnny                 1
         Debussy, Claude               1
         Baxter, Phil                  1
         Cannon, Hughie                1
         Confrey, Zez                  1
         Anthem,                       1
         Ellington, Duke               1
         Sampson, Edgar                1
         Handy, William Christopher    1
         Name: composerName, dtype: int64
```

2 RESTful Web Services

Many RESTful web services return data in XML format. Like before, we use the `requests` library in Python to issue the HTTP request. For example, the website [FloatRates](http://www.floatrates.com/) provides exchange rates between world currencies in XML format.

```
In [17]: import requests
         resp = requests.get("http://www.floatrates.com/daily/usd.xml")
         resp
```

```
Out[17]: <Response [200]>
```

The XML is stored in the `.content` attribute of the response object. We can parse this string into an `ElementTree` using the `.fromstring()` function in the `lxml` library. Recall that this returns the root tag of the XML document.

```
In [18]: etree.fromstring(resp.content)
```

```
Out[18]: <Element channel at 0x7f4a3989f848>
```

3 Exercises

Exercises 1 and 2 deal with the New York Philharmonic data set from above. These exercises are the same as the ones in the previous section, except that now you have to do them with XML.

Exercise 1. What is the most frequent start time for New York Philharmonic concerts?

```
In [19]: rows = []

concerts = programs.findall("./concertInfo")
for concert in concerts:
    row = {}
    row["Time"] = concert.find("Time").text
    row["Season"] = concert.xpath("ancestor::program//season")[0].text
    rows.append(row)

concertsdf = pd.DataFrame(rows)
concertsdf.Time.value_counts().head()
```

```
Out[19]: 8:30PM      4584
         8:00PM      4443
         3:00PM      2133
         7:30PM      2075
         2:30PM      1618
         Name: Time, dtype: int64
```

Exercise 2. How many total concerts did the New York Philharmonic perform in the 2014-15 season?

```
In [20]: len(concertsdf[concertsdf.Season == "2014-15"])
```

```
Out[20]: 217
```

In Exercises 3-4, you will work with [APIXU](#), an weather API. This API returns data in both JSON and XML formats. In these exercises, you should request the data to be returned in XML format.

Register with the website to obtain an API key. You will likely need to refer to [the API documentation here](#). If you run into unexpected errors, issue the HTTP request from your browser to make sure that the data is in the format you expect.

Exercise 3. Get the forecasted low (min) and high (max) temperatures (in Fahrenheit) for the next 7 days in San Luis Obispo. Make a graphic that displays this information.

```
In [21]: apikey = "be1c5704a83c459ea4f162211192802"
         resp = requests.get(
             "https://api.apixu.com/v1/forecast.xml?q=San Luis Obispo&days=7&key=%s" %apikey)

In [22]: wea = etree.fromstring(resp.content)
         %matplotlib inline
```

```


In [23]: rows = []

forecastdays = wea.xpath("./forecastday")
for forecastday in forecastdays:
    row = {}
    row["min_temp"] = pd.to_numeric(forecastday.xpath("./mintemp_f")[0].text)
    row["max_temp"] = pd.to_numeric(forecastday.xpath("./maxtemp_f")[0].text)
    rows.append(row)

forecastdf = pd.DataFrame(rows)
forecastdf.plot.bar()

```

Out[23]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4a1ccad940>



Exercise 4. Get the hourly wind speed (in mph) for the past 7 days. (*Note:* This will require making 7 HTTP requests to the API. Try to do it programmatically.) You should end up with $24 \times 7 = 168$ rows in your DataFrame. Make a plot of the wind speed as a function of time. What do you notice?

```

In [27]: apikey = "be1c5704a83c459ea4f162211192802"
        resp = requests.get(
            "https://api.apixu.com/v1/history.xml?q=San Luis Obispo&dt=2019-02-27&key=%s"

In [34]: root = etree.fromstring(resp.content)
        #print(etree.tostring(root, pretty_print=True).decode())

In [35]: import time
        len(root.xpath("./hour"))

```

```

rows = []
for day in range(21,28):
    date = "2019-02-%d" % day

    resp = requests.get(
        "https://api.apixu.com/v1/history.xml?q=San Luis Obispo&dt=%s&key=%s" % (date

    root = etree.fromstring(resp.content)


    for hour in root.xpath("./hour"):
        rows.append(float(hour.xpath("./wind_mph")[0].text))

    time.sleep(0.1)

pd.Series(rows).plot.line()

```

Out [35]: <matplotlib.axes._subplots.AxesSubplot at 0x7f4a1be1eb38>



```
In [ ]: hour = root.xpath("./hour")
```