

8A. Song Lyrics Generator

March 11, 2019

1 8A. Song Lyrics Generator

In this lab, you will scrape a website to get lyrics of songs by your favorite artist. Then, you will train a model called a Markov chain on these lyrics so that you can generate a song in the style of your favorite artist.

2 Question 1. Scraping Song Lyrics

Find a web site that has lyrics for several songs by your favorite artist. Scrape the lyrics into a Python list called `lyrics`, where each element of the list represents the lyrics of one song.

Tips: - Find a web page that has links to all of the songs, like [this one](#). [Note: It appears that `azlyrics.com` blocks web scraping, so you'll have to find a different lyrics web site.] Then, you can scrape this page, extract the hyperlinks, and issue new HTTP requests to each hyperlink to get each song. - Use `time.sleep()` to stagger your HTTP requests so that you do not get banned by the website for making too many requests.

```
In [1]: import requests
import time

from bs4 import BeautifulSoup

names = ["gods-plan", "hotline-bling", "hold-on-were-going-home", "one-dance",
         "in-my-feelings", "nice-for-what", "find-your-love",
         "too-good", "best-i-ever-had", "passionfruit"]

lyrics = []

for i in names:
    resp = requests.get("http://www.metrolyrics.com/{0}-lyrics-drake.html".format(i))
    time.sleep(.5)
    soup = BeautifulSoup(resp.content)

    song = soup.find_all("p", attrs={"class": "verse"})

    final_lyrics = ""
```

```

time.sleep(.25)
for i in range(len(song)):
    verse = song[i].text.strip()
    #print(verse)
    #final_lyrics = " <N> " + verse
    final_lyrics += verse

lyrics.append(final_lyrics)

In [2]: lyrics[0] = lyrics[0].replace('[Intro]', '')
lyrics[0] = lyrics[0].replace('[Chorus 1]', '')
lyrics[0] = lyrics[0].replace('[Chorus 2]', '')
lyrics[0] = lyrics[0].replace('[Verse 1]', '')
lyrics[0] = lyrics[0].replace('[Verse 2]', '')
lyrics[0] = lyrics[0].replace('[Bridge]', '')
for i in range(len(lyrics)):
    lyrics[i] = lyrics[i].replace('\n', " <N> ")

```

pickle is a Python library that serializes Python objects to disk so that you can load them in later.

```

In [3]: import pickle
        pickle.dump(lyrics, open("lyrics.pkl", "wb"))

```

3 Question 2. Unigram Markov Chain Model

You will build a Markov chain for the artist whose lyrics you scraped in Lab A. Your model will process the lyrics and store the word transitions for that artist. The transitions will be stored in a dict called chain, which maps each word to a list of “next” words.

For example, if your song was *“The Joker”* by the [Steve Miller Band](#), chain might look as follows:

```

chain = {
    "some": ["people", "call", "people"],
    "call": ["me", "me", "me"],
    "the": ["space", "gangster", "pompitous", ...],
    "me": ["the", "the", "Maurice"],
    ...
}

```

Besides words, you should include a few additional states in your Markov chain. You should have "<START>" and "<END>" states so that we can keep track of how songs are likely to begin and end. You should also include a state called "<N>" to denote line breaks so that you can keep track of where lines begin and end. It is up to you whether you want to include normalize case and strip punctuation.

So for example, for *“The Joker”*, you would add the following to your chain:

```

chain = {

```

```

"<START>": ["Some", ...],
"Some": ["people", ...],
"people": ["call", ...],
"call": ["me", ...],
"me": ["the", ...],
"the": ["space", ...],
"space": ["cowboy,", ...],
"cowboy,": ["yeah", ...],
"yeah": ["<N>", ...],
"<N>": ["Some", ..., "Come"],
...,
"Come": ["on", ...],
"on": ["baby", ...],
"baby": ["and", ...],
"and": ["I'll", ...],
"I'll": ["show", ...],
"show": ["you", ...],
"you": ["a", ...],
"a": ["good", ...],
"good": ["time", ...],
"time": ["<END>", ...],
}

```

Your chain will be trained on not just one song, but by all songs by your artist.

```

In [4]: def train_markov_chain(lyrics):
        """
        Args:
            - lyrics: a list of strings, where each string represents
                      the lyrics of one song by an artist.

        Returns:
            A dict that maps a single word ("unigram") to a list of
            words that follow that word, representing the Markov
            chain trained on the lyrics.
        """
        chain = {"<START>" : []}

        for lyric in lyrics:
            count = 0
            #print(lyric)
            words = lyric.split()
            #print(words)
            chain["<START>"].append(words[0])
            for i in range(len(words)):
                if i < len(words) - 1:
                    if words[i] in chain:
                        chain[words[i]].append(words[i+1])

```

```

        else:
            chain[words[i]] = [words[i+1]]
    else:
        if words[i] in chain:
            chain[words[i]].append("<END>")
            #chain[words[i]].append("END")
        else:
            chain[words[i]] = ["<END>"]
            #chain[words[i]] = "END"

    return chain

```

In [5]: *# Load the pickled lyrics object that you created in Lab A.*

```
import pickle
```

```
lyrics = pickle.load(open("lyrics.pkl", "rb"))
```

Call the function you wrote above.

```
chain = train_markov_chain(lyrics)
```

What words tend to start a song (i.e., what words follow the <START> tag?)

```
print(chain["<START>"])
```

What words tend to begin a line (i.e., what words follow the line break tag?)

```
print(chain["<N>"][:20])
```

```
['<N>', 'You', 'I', 'Baby,', 'Trap,', 'I', 'I'm', 'Oh', 'You', 'Hold']
```

```
['Yeah', 'They', 'I', 'Tryna', "Don't", 'You', 'I', 'Guess', 'Hope', 'They', "God's", 'I', 'I'
```

Now, let's generate new lyrics using the Markov chain you constructed above. To do this, we'll begin at the "<START>" state and randomly sample a word from the list of words that follow "<START>". Then, at each step, we'll randomly sample the next word from the list of words that followed each current word. We will continue this process until we sample the "<END>" state. This will give us the complete lyrics of a randomly generated song!

You may find the `random.choice()` function helpful for this question.

In [6]: `import random`

```
def generate_new_lyrics(chain):
```

```
    """
```

Args:

```
    - chain: a dict representing the Markov chain,
              such as one generated by generate_new_lyrics()
```

Returns:

```
    A string representing the randomly generated song.
```

```
    """
```

```
    # a list for storing the generated words
```

```

words = []
# generate the first word
words.append(random.choice(chain("<START>")))

# YOUR CODE HERE
done = False
while not done:
    words.append(random.choice(chain(words[len(words) - 1])))
    if words[len(words) - 1] == "<END>":
        done = True

# join the words together into a string with line breaks
lyrics = " ".join(words[:-1])
return "\n".join(lyrics.split("<N>"))

```

```
In [7]: print(generate_new_lyrics(chain))
```

```

I just being right
It feels like you always stay at me
Give to ask you my hands around me
Just hold onI keep letting you need a patient in your lovin'
Wonder if you're okay
Last night I need that hotline bling
Why won't mention
You're everything
I hope that you the fuckin' best
Got a realization
Said you'd die for granted
That they wishin'
And when that ass back up
Just hold on Southside G, yu
That can I can't get the rumors and they wishin'
She say, "Do you the friendship
Are you see if I ever leave
First, last slice
I bet if this forever
I'll be real one, in your heart
KB, do you know you been peepin' what they wishin'
Hangin' with you
First, last slice
Just make it peaceful is a one
I get rightI promise that ass
You take itI'm too good to call me
'Cause if you're a nigga hit it, yupKiki, do it
You ain't even know it taking a G.O.D., yeah, wait a real big
All up and I promise that I need my love my bed and me, care for me!
And I'm more time I can only mean one thing every single time I left your lovin'

```

You said you'd die for granted
Watch the fact that...I'm too good girl and I won't mention
Give to get a hold on my eyes on me to me!
There for me!
Wonder if I better find your expectations
No, I'm away
You the new me down pon it from the only mean one that shit
You said you'd die for me!
Hey, hey, hey
Cock up on my heart
And when friends always touching road
Then nothing's gonna tear us apart
Going places where do you
Every time 'fore I need ya
I just leave
You hear but you everything
I get some Henny
You act so long
Doing overtime for me! (A song
I want ya, and wishin' and they running out commitment for me!
I know you
Back up, back up on we're going home
Hey, hey, hey
Said you'd die for you used to these niggas
I can't end like I ever had Best I think there's never alone (things alone)
That can spend whatever on me
I can't miss
I'll be forgotten
They wishin'
I give to call me is costing me
I need ya
I hope you know shorty and wishin' and I ever had
Are you ridin' say you'll never run away
You don't get em gassed up
Have a hold on we're going home (home) I got some brothers that hotline bling
I do you my wallet
Last night, I don't understand it Years go
Higher powers taking a one dance
You the safe
Gotta make your mothafuckin' roll on my hand
Back up, back up the same thing
Young money ain't ever had
It's a Hennessy in
Feelin' for me, 'cause I get over backwards for it
Passive with me feel like you
It's hard at buildin' trust from miles away
You take my momma, I'm way too good to you are...
I think I just gotta be there for a hold on Southside G, ay, road, ay

High school pics, you don't understand it like the fuckin' best
 'Cause we're going home (home)I got to you andI need my love for me, care for me, cry for me?
 Why won't you ain't ever had
 Passive with me when its stopping
 And when that ass back in so accustomed to the key under the fuckin' bra strap pop
 There for my love and wishin' and me, we gon' live for now
 Got a distance
 They wishin' and wishin' and I just be over you the same thing every single show she out regul
 I ever wanted
 You said you'd be over backwards for me?!Gotta hit her up on me
 I'm down
 You act so official
 You used to somebody
 I thought you know alot of pages in my patience
 I can't blame you, no, no me things
 Oti, oti, there's nothing left to you the city, you know it for me!
 They gon' live it tattted on Southside G, ay, ay
 Then nothing's gonna tear us apartToo many times
 You know it up a mission
 Mi wi give to my love me
 I don't see the one thing to you to me!
 There for it up on
 Wonder if you're down pon it tattted on it like it
 B-bring that hotline bling
 I had
 I better find your slot 'til a mission
 Then nothing's gonna tear us apartToo many times I've been inside, know too and I guess being
 Thinkin' my love
 I swear I give to the fuckin' best
 I give it real with me
 Said you'd be there as you need a blank disc
 I'll be no lover in here tonight
 And I'm down
 Why you

4 Question 3. Bigram Markov Chain Model

Now you'll build a more complex Markov chain that uses the last *two* words (or bigram) to predict the next word. Now your dict chain should map a *tuple* of words to a list of words that appear after it.

As before, you should also include tags that indicate the beginning and end of a song, as well as line breaks. That is, a tuple might contain tags like "<START>", "<END>", and "<N>", in addition to regular words. So for example, for ["The Joker"](#), you would add the following to your chain:

```
chain = {
    (None, "<START>"): ["Some", ...],
```

```

("<START>", "Some"): ["people", ...],
("Some", "people"): ["call", ...],
("people", "call"): ["me", ...],
("call", "me"): ["the", ...],
("me", "the"): ["space", ...],
("the", "space"): ["cowboy,", ...],
("space", "cowboy,"): ["yeah", ...],
("cowboy,", "yeah"): ["<N>", ...],
("yeah", "<N>"): ["Some", ...],
("time", "<N>"): ["Come"],
...,
("<N>", "Come"): ["on", ...],
("Come", "on"): ["baby", ...],
("on", "baby"): ["and", ...],
("baby", "and"): ["I'll", ...],
("and", "I'll"): ["show", ...],
("I'll", "show"): ["you", ...],
("show", "you"): ["a", ...],
("you", "a"): ["good", ...],
("a", "good"): ["time", ...],
("good", "time"): ["<END>", ...],
}

```

```

In [8]: def train_markov_chain(lyrics):
        """
        Args:
            - lyrics: a list of strings, where each string represents
                      the lyrics of one song by an artist.

        Returns:
            A dict that maps a tuple of 2 words ("bigram") to a list of
            words that follow that bigram, representing the Markov
            chain trained on the lyrics.
        """
        chain = {(None, "<START>"): []}
        for lyric in lyrics:
            # YOUR CODE HERE
            count = 0
            #print(lyric)
            words = lyric.split()
            #print(words)
            chain[(None, "<START>")].append(words[0])
            chain[("<START>", words[0])] = [words[1]]
            for i in range(len(words)-1):
                if i < len(words) - 2:
                    if words[i] in chain and words[i+1] in chain:
                        chain[(words[i], words[i+1])].append(words[i+2])
                    else:

```



```

        chain[(words[i], words[i+1])] = [words[i+2]]
    else:
        if words[i] in chain and words[i+1] in chain:
            chain[(words[i], words[i+1])].append("<END>")
            #chain[words[i]].append("END")
        else:
            chain[(words[i], words[i+1])] = ["<END>"]
            #chain[words[i]] = "END"

    return chain

In [9]: # Load the pickled lyrics object that you created in Lab A.
import pickle
lyrics = pickle.load(open("lyrics.pkl", "rb"))

# Call the function you wrote above.
chain = train_markov_chain(lyrics)

# What words tend to start a song (i.e., what words follow the <START> tag?)
print(chain[(None, "<START>")])

['<N>', 'You', 'I', 'Baby,', 'Trap,', 'I', "I'm", 'Oh', 'You', 'Hold']

```

Now, let's generate new lyrics using the Markov chain you constructed above. To do this, we'll begin at the (None, "<START>") state and randomly sample a word from the list of words that follow this bigram. Then, at each step, we'll randomly sample the next word from the list of words that followed the current bigram (i.e., the last two words). We will continue this process until we sample the "<END>" state. This will give us the complete lyrics of a randomly generated song!

```

In [10]: import random

def generate_new_lyrics(chain):
    """
    Args:
        - chain: a dict representing the Markov chain,
                  such as one generated by generate_new_lyrics()

    Returns:
        A string representing the randomly generated song.
    """

    # a list for storing the generated words
    words = []
    # generate the first word
    words.append(random.choice(chain[(None, "<START>")]))
    words.append(random.choice(chain[("<START>", words[0])]))

```

```

# YOUR CODE HERE
done = False
while not done:
    words.append(random.choice(chain[(words[len(words)-2], words[len(words)-1])]))

    #if words[len(words) - 1] == "the":
    if words[len(words) - 1] == "<END>":
        done = True

# join the words together into a string with line breaks
lyrics = " ".join(words[:-1])
return "\n".join(lyrics.split("<N>"))

```

```
In [11]: print(generate_new_lyrics(chain))
```

```

I'm more than it seems
Sacrificing things
And I say you the f...
Uh
Oh yeah
Tension between us just like picket fences
You got issues that I won't mention for now
They keep fallin' apartPassionate from miles away
Passive with the things you say
Passin' up on my old ways
I can't blame you, no, noTrying to think of the right thing to say

```

5 Analysis

Compare the quality of the lyrics generated by the unigram model (in Lab B) and the bigram model (in Lab C). Which model seems to generate more reasonable lyrics? Can you explain why? What do you see as the advantages and disadvantages of each model?

The lyrics generated from the unigram model are not as fluid as the bigram model. The lyrics make more sense in the bigram model because it is using pairs of words so then there will be more unique keys. Then for each key, the amount of values to sample from will be smaller. Therefore, the bigram will produce lyrics that read easier. However, one disadvantage to this is that we lose the “random” effect, and predicting the lyrics might be easier if the chain starts off with unique tuples of keys. This in turn leads to actual repeated lyrics from the artist’s songs. For example, in the bigram above, lines 8 and 9 are word for word with Drake’s song “Passionfruit”.

6 Submission Instructions

Once you are finished, follow these steps:

1. Restart the kernel and re-run this notebook from beginning to end by going to Kernel > Restart Kernel and Run All Cells.

2. If this process stops halfway through, that means there was an error. Correct the error and repeat Step 1 until the notebook runs from beginning to end.
3. Double check that there is a number next to each code cell and that these numbers are in order.

Then, submit your lab as follows:

1. Go to File > Export Notebook As > PDF.
2. Double check that the entire notebook, from beginning to end, is in this PDF file. (If the notebook is cut off, try first exporting the notebook to HTML and printing to PDF.)
3. Upload the PDF [to PolyLearn](#).