

# Exam2

March 7, 2019

## 1 Exam 2

In this exam, you will request data from the [Zomato API](#) and then train machine learning models on the data that you gather. You should have already registered an API key with Zomato, which will entitle you to 1000 API requests per day. This exam can be completed with fewer than 200 API requests.

```
In [1]: %matplotlib inline

import numpy as np
import pandas as pd
import requests

apikey = "49f040db06fd803e2a8465d4a241daa3"
```

## 2 Question 1 (4 points)

Determine the Zomato city ID for Chicago, IL. Then, use this city ID to create a DataFrame containing all of the cuisines in Chicago, along with their IDs. Display this DataFrame.

```
In [2]: url = ("https://developers.zomato.com/api/v2.1/"
              "location_details?entity_id=292&entity_type=city")

resp = requests.get(url, headers={"user-key": apikey})
chicago = resp.json()

In [3]: import json
        from pandas.io.json import json_normalize

chicago_df = json_normalize(chicago["bestRatedRestaurant"])
chicago_df.head()

Out[3]:
```

	restaurant.res_id	restaurant.apikey \
0	16736014	49f040db06fd803e2a8465d4a241daa3
1	16737455	49f040db06fd803e2a8465d4a241daa3
2	16752484	49f040db06fd803e2a8465d4a241daa3
3	16734364	49f040db06fd803e2a8465d4a241daa3

```

4          16753200  49f040db06fd803e2a8465d4a241daa3

restaurant.average_cost_for_two restaurant.book_again_url \
0          40
1          30
2          35
3          90
4          85

restaurant.book_form_web_view_url restaurant.cuisines \
0          Pizza
1          American, Italian, Burger, Sandwich
2          Mexican
3          Brazilian, Steak
4          New American

restaurant.currency restaurant.deeplink \
0          $ zomato://restaurant/16736014
1          $ zomato://restaurant/16737455
2          $ zomato://restaurant/16752484
3          $ zomato://restaurant/16734364
4          $ zomato://restaurant/16753200

restaurant.events_url \
0 https://www.zomato.com/chicago/lou-malnatis-pi...
1 https://www.zomato.com/chicago/portillos-hot-d...
2 https://www.zomato.com/chicago/xoco-river-nort...
3 https://www.zomato.com/chicago/fogo-de-chao-br...
4 https://www.zomato.com/chicago/girl-the-goat-w...

restaurant.featured_image \
0 https://b.zmtcdn.com/data/res_imagery/16736012...
1 https://b.zmtcdn.com/data/res_imagery/16737455...
2 https://b.zmtcdn.com/data/res_imagery/16752484...
3 https://b.zmtcdn.com/data/res_imagery/16734364...
4 https://b.zmtcdn.com/data/res_imagery/16753200...

... \
0 ...
1 ...
2 ...
3 ...
4 ...

restaurant.photos_url restaurant.price_range \
0 https://www.zomato.com/chicago/lou-malnatis-pi... 2
1 https://www.zomato.com/chicago/portillos-hot-d... 2
2 https://www.zomato.com/chicago/xoco-river-nort... 2

```

3	<a href="https://www.zomato.com/chicago/fogo-de-chao-br...">https://www.zomato.com/chicago/fogo-de-chao-br...</a>	4
4	<a href="https://www.zomato.com/chicago/girl-the-goat-w...">https://www.zomato.com/chicago/girl-the-goat-w...</a>	4

restaurant.switch_to_order_menu \	
0	0
1	0
2	0
3	0
4	0

restaurant.thumb \	
0	<a href="https://b.zmtcdn.com/data/res_imagery/16736012...">https://b.zmtcdn.com/data/res_imagery/16736012...</a>
1	<a href="https://b.zmtcdn.com/data/res_imagery/16737455...">https://b.zmtcdn.com/data/res_imagery/16737455...</a>
2	<a href="https://b.zmtcdn.com/data/res_imagery/16752484...">https://b.zmtcdn.com/data/res_imagery/16752484...</a>
3	<a href="https://b.zmtcdn.com/data/res_imagery/16734364...">https://b.zmtcdn.com/data/res_imagery/16734364...</a>
4	<a href="https://b.zmtcdn.com/data/res_imagery/16753200...">https://b.zmtcdn.com/data/res_imagery/16753200...</a>

restaurant.url \	
0	<a href="https://www.zomato.com/chicago/lou-malnatis-pi...">https://www.zomato.com/chicago/lou-malnatis-pi...</a>
1	<a href="https://www.zomato.com/chicago/portillos-hot-d...">https://www.zomato.com/chicago/portillos-hot-d...</a>
2	<a href="https://www.zomato.com/chicago/xoco-river-nort...">https://www.zomato.com/chicago/xoco-river-nort...</a>
3	<a href="https://www.zomato.com/chicago/fogo-de-chao-br...">https://www.zomato.com/chicago/fogo-de-chao-br...</a>
4	<a href="https://www.zomato.com/chicago/girl-the-goat-w...">https://www.zomato.com/chicago/girl-the-goat-w...</a>

restaurant.user_rating.aggregate_rating \	
0	4.7
1	4.9
2	4.9
3	4.6
4	4.8

restaurant.user_rating.has_fake_reviews \	
0	0
1	0
2	0
3	0
4	0

restaurant.user_rating.rating_color	restaurant.user_rating.rating_text \
0	3F7E00      Excellent
1	3F7E00      Excellent
2	3F7E00      Excellent
3	3F7E00      Excellent
4	3F7E00      Excellent

restaurant.user_rating.votes	
0	789
1	1023

2	737
3	601
4	743

[5 rows x 43 columns]

### 3 Question 2 (4 points)

Get the top 40 restaurants (sorted in desc order by rating) in Chicago for the following cuisines:

- Chinese
- Italian
- Mexican

Store the 120 results in a single DataFrame with the following columns:

- res\_id: the ID for the Zomato restaurant
- name
- all of the location features (i.e., address, locality, latitude, longitude, zipcode, etc.)
- all of the user\_rating features (i.e., aggregate\_rating, votes, etc.)

Display this DataFrame.

```
In [4]: url = ("https://developers.zomato.com/api/v2.1/search?entity_id=292&entity"
              "_type=city&start=0&count=40&cuisines=25&sort=rating&order=desc"
              )
resp = requests.get(url, headers={"user-key": apikey})
chicago_chinese_part1 = resp.json()

chicago_chinese_df_part1 = json_normalize(chicago_chinese_part1["restaurants"])

In [5]: url = ("https://developers.zomato.com/api/v2.1/search?"
              "entity_id=292&entity_type=city&start=20&cuisines=25&sort=rating&order=desc")

resp = requests.get(url, headers={"user-key": apikey})
chicago_chinese_part2 = resp.json()

chicago_chinese_df_part2 = json_normalize(chicago_chinese_part2["restaurants"])

In [6]: url = ("https://developers.zomato.com/api/v2.1/search?"
              "entity_id=292&entity_type=city&start=0&cuisines=55&sort=rating&order=desc")

resp = requests.get(url, headers={"user-key": apikey})
chicago_italian_part1 = resp.json()

chicago_italian_part1_df = json_normalize(chicago_italian_part1["restaurants"])
```

```

In [7]: url = ("https://developers.zomato.com/api/v2.1/search?entity_id"
              "=292&entity_type=city&start=20&cuisines=55&sort=rating&order=desc")

resp = requests.get(url, headers={"user-key": apikey})
chicago_italian_part2 = resp.json()

chicago_italian_part2_df = json_normalize(chicago_italian_part2["restaurants"])

In [8]: url = ("https://developers.zomato.com/api/v2.1/search?entity_id="
              "292&entity_type=city&start=0&cuisines=73&sort=rating&order=desc")

resp = requests.get(url, headers={"user-key": apikey})
chicago_mexican_part1 = resp.json()

chicago_mexican_part1_df = json_normalize(chicago_mexican_part1["restaurants"])

In [9]: url = ("https://developers.zomato.com/api/v2.1/search?entity_id="
              "292&entity_type=city&start=20&cuisines=73&sort=rating&order=desc")

resp = requests.get(url, headers={"user-key": apikey})
chicago_mexican_part2 = resp.json()

chicago_mexican_part2_df = json_normalize(chicago_mexican_part2["restaurants"])

In [10]: chicago_chinese_df_part1["my_cuisine"] = "Chinese"
chicago_chinese_df_part2["my_cuisine"] = "Chinese"
chicago_italian_part1_df["my_cuisine"] = "Italian"
chicago_italian_part2_df["my_cuisine"] = "Italian"
chicago_mexican_part1_df["my_cuisine"] = "Mexican"
chicago_mexican_part2_df["my_cuisine"] = "Mexican"

all_dfs = [chicago_chinese_df_part1, chicago_chinese_df_part2,
            chicago_italian_part1_df, chicago_italian_part2_df,
            chicago_mexican_part1_df, chicago_mexican_part2_df]

chicago_restaurants = pd.concat(all_dfs)

In [11]: my_index = range(0,120)

chicago_restaurants["my_index"] = my_index

chicago_restaurants = chicago_restaurants.set_index(chicago_restaurants["my_index"])

In [12]: final_chicago_df = chicago_restaurants[["restaurant.R.res_id", "my_cuisine",
            "restaurant.location.address", "restaurant.location.city",
            "restaurant.location.city_id", "restaurant.location.country_id",
            "restaurant.location.latitude", "restaurant.location.locality",
            "restaurant.location.locality_verbose",
            "restaurant.location.longitude", "restaurant.location.zipcode",

```

```

"restaurant.name", "restaurant.user_rating.aggregate_rating",
"restaurant.user_rating.rating_color",
"restaurant.user_rating.rating_text",
"restaurant.user_rating.votes"]]]

```

```

final_chicago_df.head()

```

```

Out[12]:          restaurant.R.res_id my_cuisine \

```

```

my_index
0          16735363    Chinese
1          16749821    Chinese
2          16743712    Chinese
3          16747211    Chinese
4          16751533    Chinese

```

```

          restaurant.location.address \

```

```

my_index
0                    521 Davis Street 60201
1                    200 E Golf Road 60173
2          9560 S. Kedzie Avenue, Evergreen Park 60805
3          1003 W. Ogden Avenue, Suite B, Naperville 60563
4          537 Green Bay Road, Wilmette 60091

```

```

          restaurant.location.city restaurant.location.city_id \

```

```

my_index
0          Chicago          292
1          Chicago          292
2          Chicago          292
3          Chicago          292
4          Chicago          292

```

```

          restaurant.location.country_id restaurant.location.latitude \

```

```

my_index
0          216          42.0460805556
1          216          42.0507000000
2          216          41.7197000000
3          216          41.7848777778
4          216          42.0748055556

```

```

          restaurant.location.locality restaurant.location.locality_verbose \

```

```

my_index
0          Evanston          Evanston, Chicago
1          Schaumburg          Schaumburg, Chicago
2          Evergreen Park          Evergreen Park, Chicago
3          Naperville          Naperville, Chicago
4          Wilmette          Wilmette, Chicago

```

```

          restaurant.location.longitude restaurant.location.zipcode \

```

```

my_index
0          -87.6790305556          60201
1          -88.0742638889          60173
2          -87.7020000000          60805
3          -88.1680916667          60563
4          -87.7077638889          60091

          restaurant.name restaurant.user_rating.aggregate_rating \
my_index
0          Joy Yee's Noodles          4.5
1          Yu's Mandarin Restaurant          4.5
2          Chi Tung Restaurant          4.5
3          Chinese Kitchen          4.4
4          Tsing Tao Mandarin Chinese          4.4

          restaurant.user_rating.rating_color \
my_index
0          3F7E00
1          3F7E00
2          3F7E00
3          5BA829
4          5BA829

          restaurant.user_rating.rating_text restaurant.user_rating.votes
my_index
0          Excellent          161
1          Excellent          158
2          Excellent          115
3          Very Good          31
4          Very Good          22

```

#### 4 Question 3 (4 points)

For each of the 120 restaurants that you identified in Question 2, get the 5 most recent reviews. Store all of the reviews in a DataFrame with the following columns:

- rating
- review\_text
- all of the user features (such as name, foodie\_level, foodie\_level\_num)

Display this DataFrame.

```

In [13]: import time

all_reviews_df = pd.DataFrame()

for res_id in final_chicago_df["restaurant.R.res_id"]:
    url = "https://developers.zomato.com/api/v2.1/reviews?res_id={}".format(res_id)

```

```

resp = requests.get(url, headers={"user-key": apikey})
reviews = resp.json()
reviews_df = json_normalize(reviews["user_reviews"])
all_reviews_df = all_reviews_df.append(reviews_df)

time.sleep(.5)

```

/opt/conda/lib/python3.6/site-packages/pandas/core/frame.py:6211: FutureWarning: Sorting because  
of pandas will change to not sort by default.

To accept the future behavior, pass 'sort=False'.

To retain the current behavior and silence the warning, pass 'sort=True'.

```
sort=sort)
```

```

In [14]: my_index = range(0,578)
         all_reviews_df["my_index"] = my_index

         all_reviews_df = all_reviews_df.set_index(all_reviews_df["my_index"])

```

```
In [15]: all_reviews_df.head()
```

```

Out[15]:
          review.comments_count  review.id  review.likes  review.rating \
my_index
0                0  32160086                0                5.0
1                0  26207391                1                5.0
2                0  24891125                0                4.0
3                0  24622388                1                3.5
4                0  16100548                0                0.0

```

```

          review.rating_color  review.rating_text \
my_index
0                305D02          Insane!
1                305D02          Insane!
2                5BA829          Great!
3                9ACD32      Good Enough
4                CBCBC8      Not rated

```

```

          review.review_text \
my_index
0      Lively ambiance, but its all about the food. T...
1      It was very good food a customer service Has ...
2      Yummy yummy yummy! This place is a gold mine. ...
3      Huge portions. good food, and they actually ma...
4      Stay away from this rapidly declining restaura...

```

```
review.review_time_friendly  review.timestamp \
```



my_index		
0	Oct 23, 2017	1508781075
1	Feb 07, 2016	1454800032
2	Sep 11, 2015	1441916515
3	Aug 15, 2015	1439587312
4	Nov 13, 2013	1384366238

	review.user.foodie_color	review.user.foodie_level	\
my_index			
0	ffd35d	Foodie	
1	ffd35d	Foodie	
2	ffae4f	Big Foodie	
3	f58552	Super Foodie	
4	ffd35d	Foodie	

	review.user.foodie_level_num	review.user.name	\
my_index			
0	3	Jason Jobe	
1	2	Tarek Anthony	
2	6	Jalpa T	
3	8	Prapti	
4	1	Mtnsguy	

	review.user.profile_deeplink	\
my_index		
0	zomato://u/30860504	
1	zomato://u/33462700	
2	zomato://u/30483347	
3	zomato://u/1510791	
4	zomato://u/23884812	

	review.user.profile_image	\
my_index		
0	https://b.zmtcdn.com/data/user_profile_picture...	
1	https://b.zmtcdn.com/data/user_profile_picture...	
2	https://b.zmtcdn.com/data/user_profile_picture...	
3	https://b.zmtcdn.com/data/user_profile_picture...	
4	https://b.zmtcdn.com/images/user_avatars/mug_2...	

	review.user.profile_url	\
my_index		
0	https://www.zomato.com/users/jason-jobe-308605...	
1	https://www.zomato.com/users/tarek-anthony-334...	
2	https://www.zomato.com/samosapop?utm_source=ap...	
3	https://www.zomato.com/Praptisahni?utm_source=...	
4	https://www.zomato.com/users/mtnsguy-23884812?...	

	review.user.zomato_handle	my_index
--	---------------------------	----------

my_index		
0	NaN	0
1	NaN	1
2	samosapop	2
3	Praptisahni	3
4	NaN	4

## 5 Question 4 (6 points)

Let's use the restaurants data that you obtained in Question 2 to train a machine learning model to predict whether a restaurant serves Chinese, Italian, or Mexican cuisine, given just the latitude and the longitude of the restaurant. (For a 1 point penalty, you can download `restaurants.csv` from [PolyLearn](#), upload it to the current directory, and read in the file.)

Train a  $k$ -nearest neighbors model. Determine the optimal value of  $k$ . Calculate an estimate of the test precision and recall of your final model. Interpret these values in the context of this application.

```
In [16]: from sklearn.preprocessing import StandardScaler
         from sklearn.neighbors import KNeighborsClassifier
         from sklearn.metrics import precision_score, recall_score
         from sklearn.pipeline import Pipeline
         from sklearn.model_selection import cross_val_score
         from sklearn.metrics import accuracy_score
         from sklearn.metrics import f1_score

In [17]: import warnings
         warnings.filterwarnings('ignore')

         features = ["restaurant.location.latitude", "restaurant.location.longitude"]

         X_train = final_chicago_df[features].astype('float64')
         y_train = final_chicago_df["my_cuisine"]

         is_Chinese_train = (y_train == "Chinese")

         scaler = StandardScaler()
         scaler.fit(X_train)
         X_train_sc = scaler.transform(X_train)

         best_k = []
         best_cv = []
         best_accuracy = []
         best_recall = []
         best_precision = []
         best_score = 0

         for k in list(range(1,31)):
             model = KNeighborsClassifier(n_neighbors=k)
```

```

model.fit(X_train_sc, y_train)

pipeline = Pipeline([
    ("scaler", scaler),
    ("model", model)
])

for cross in list(range(2,11)):
    f1 = cross_val_score(pipeline, X_train, is_Chinese_train,
                        cv=cross, scoring="f1").mean()

    accuracy = cross_val_score(pipeline, X_train, y_train,
                              cv=cross, scoring="accuracy").mean()

    recall = cross_val_score(pipeline, X_train, is_Chinese_train,
                             cv=cross, scoring="recall").mean()

    precision = cross_val_score(pipeline, X_train, is_Chinese_train,
                                cv=cross, scoring="precision").mean()

    if f1 > best_score:
        best_score = f1
        best_k.append(k)
        best_cv.append(cross)
        best_accuracy.append(accuracy)
        best_recall.append(recall)
        best_precision.append(precision)

print("Neighbors: ", best_k.pop())
print("Cross: ", best_cv.pop())
print("F1: ", best_score)
print("Accuracy: ", best_accuracy.pop())
print("Recall: ", best_recall.pop())
print("Precision: ", best_precision.pop())

```

```

Neighbors:  3
Cross:  2
F1:  0.49444444444444
Accuracy:  0.425
Recall:  0.475
Precision:  0.51875

```

I have found the optimal value of k to be 3 neighbors. I trained models ranging from 1 neighbor all the way up to 30 neighbors. For each of these models, I tested cross validation folds ranging from 2 to 10. In order to reach this conclusion, I used the F1 score to compare which model was better. If the current model was better than any previous model, I saved the current model as the “best” and continued to compare the following models based on the F1 score.

## 6 Question 5 (6 points)

Let's use the reviews data that you obtained in Question 3 to train a machine learning model to predict a rating, given just the `review_text`. (For a 1 point penalty, you can download `reviews.csv` from [PolyLearn](#), upload it to the current directory, and read in the file.)

You will have to first convert the text of the review into quantitative features. Instead of including every word that appears, it is usually better to restrict to words that appear at least  $m$  times, where  $m$  is a hyperparameter. Plot the training and the test RMSE of a 10-nearest neighbors model as a function of this hyperparameter  $m$ . What value of  $m$  is optimal? What is the test RMSE of this optimal model? Interpret the test RMSE in the context of this application.

**Hint:** The hyperparameter  $m$  corresponds to the `min_df=` argument of `CountVectorizer` and `TfidfVectorizer` in *scikit-learn*.

```
In [18]: all_reviews_df["review.review_text"] =
         all_reviews_df["review.review_text"].fillna("None")

In [19]: from sklearn.feature_extraction.text import TfidfVectorizer
         from sklearn.neighbors import KNeighborsRegressor

model1 = KNeighborsRegressor(n_neighbors=10)
def test_error(m):
    m = m/10
    vec = TfidfVectorizer(norm=None, min_df=m) #, min_df=.5)
    vec.fit(all_reviews_df["review.review_text"])
    tf_idf_sparse = vec.transform(all_reviews_df["review.review_text"])
    tfidf = pd.DataFrame(tf_idf_sparse.toarray())

    X_train = tfidf
    y_train = all_reviews_df["review.rating"]

    scaler = StandardScaler()
    scaler.fit(X_train)
    X_train_sc = scaler.transform(X_train)

    model1 = KNeighborsRegressor(n_neighbors=10)
    model1.fit(X_train_sc, y_train)

    pipeline = Pipeline([
        ("scaler", scaler),
        ("model", model1)
    ])

    rmse = cross_val_score(pipeline, X_train, y_train,
                           cv=5, scoring="neg_mean_squared_error").mean()

    return np.sqrt(np.mean(-rmse))
```

```

model2 = KNeighborsRegressor(n_neighbors=10)
def train_error(m):
    m = m/10
    vec = TfidfVectorizer(norm=None, min_df=m) #, min_df=.5)
    vec.fit(all_reviews_df["review.review_text"])
    tf_idf_sparse = vec.transform(all_reviews_df["review.review_text"])
    tfidf = pd.DataFrame(tf_idf_sparse.toarray())

    X_train = tfidf
    y_train = all_reviews_df["review.rating"]

    scaler = StandardScaler()
    scaler.fit(X_train)
    X_train_sc = scaler.transform(X_train)

    model2 = KNeighborsRegressor(n_neighbors=10)
    model2.fit(X_train_sc, y_train)

    pipeline = Pipeline([
        ("scaler", scaler),
        ("model", model2)
    ])

    rmse = cross_val_score(pipeline, X_train, y_train,
                           scoring="neg_mean_squared_error").mean()

    return np.sqrt(np.mean(-rmse))

```

```

ms = pd.Series(range(0, 7, 1))
ms.index = range(0, 7, 1)
test_error = ms.apply(test_error)
train_error = ms.apply(train_error)

```

```


In [20]: test_error.plot(label="Test Error", legend=True)
train_error.plot(label="Train Error", legend=True)
test_error.sort_values()

```

```

Out[20]: 3    2.130146
         1    2.146974
         6    2.147989
         2    2.173584
         4    2.176793
         5    2.176793
         0    2.915951
dtype: float64

```



I have the optimal value of “m” to be 3 which is technically .3 as a parameter in the `min_df` argument. The rmse from using 3 as my m is 2.130146. RMSE is the average standard deviation of each data point is from its predicted point, so in context, the rmse is measuring the average distance between my model’s prediction for the review rating that each customer gave and the review rating the customer actually gave in the review.

I obtained this answer by creating 2 functions that essentially do the same thing - `train_error` does not use cross\_validation folds, while `test_error` uses 5 folds (chosen arbitrarily). I first used `TfidfVectorizer` to turn the reviews into numerical data to work with for each review. From here, it was setting the training data, scaling, making and fitting the model using pipeline from scikit learn.

## 7 Question 6 (6 points)

Let’s use the reviews data to train a machine learning model to predict the rating, given just the `foodie_level_num` of the user. Fit an 80-nearest neighbors model to predict rating from the `foodie_level_num`. Make a scatterplot showing the two variables, and add a curve to this scatterplot that shows the predicted rating as a function of `foodie_level_num`. What is the test RMSE of this model?

Then, combine this model with your (optimal) model from Question 5 into an ensemble model. How does the test RMSE of the ensemble model compare to the test RMSE of each individual model?

**Hint:** Feel free to borrow the `RegressionEnsembler` [code that I wrote](#). However, it will not work out of the box because the two models you are trying to combine use different variables as input. So if you use my `RegressionEnsembler`, you will have to adapt it to make it work for this problem.

```

In [21]: features = ["review.user.foodie_level_num"]

X_train = all_reviews_df[features]
y_train = all_reviews_df["review.rating"]

scaler = StandardScaler()
scaler.fit(X_train)
X_train_sc = scaler.transform(X_train)

model3 = KNeighborsRegressor(n_neighbors=80)
model3.fit(X_train_sc, y_train)

pipeline = Pipeline([
    ("scaler", scaler),
    ("model", model3)
])

score = cross_val_score(pipeline, X_train, y_train,
                        cv=8, scoring="neg_mean_squared_error").mean()

rmse = np.sqrt(np.mean(-score))
rmse

Out[21]: 1.4096946569859183

In [22]: def foodie_error(X_new):
        return pd.Series(model3.predict(X_new))

In [23]: X_new = pd.DataFrame()
X_new["Foodie Rating"] = np.arange(0, 14, 1)
X_new

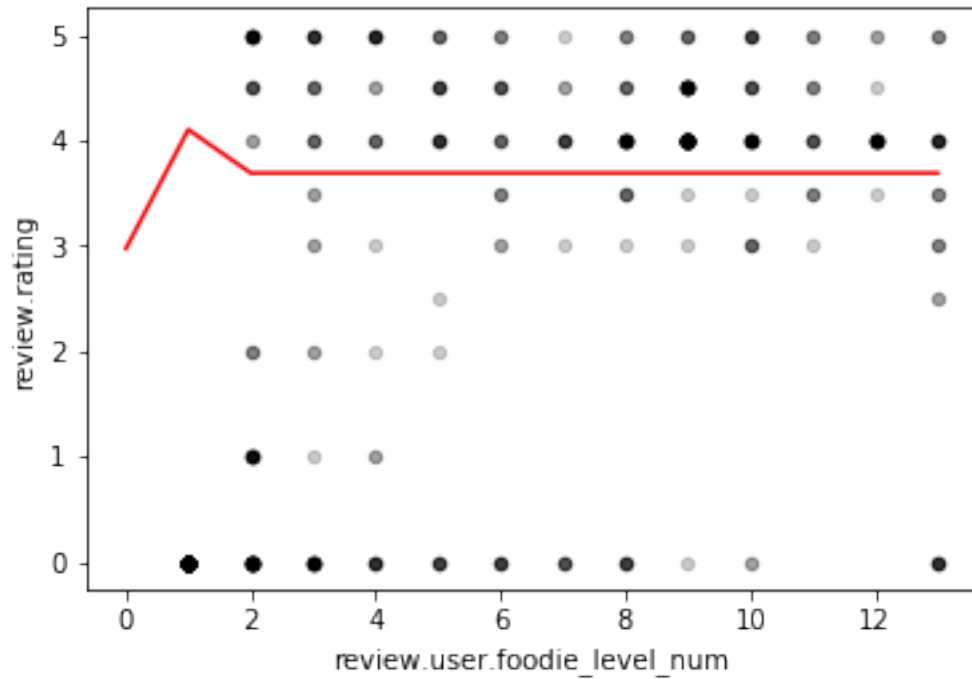
y_new_pred = foodie_error(X_new)

In [24]: all_reviews_df["review.rating"] =
        all_reviews_df["review.rating"].astype(float)

all_reviews_df.plot.scatter("review.user.foodie_level_num",
                            "review.rating", color="black", alpha=.2)
y_new_pred.index = X_new
y_new_pred.plot.line(color="red")

Out[24]: <matplotlib.axes._subplots.AxesSubplot at 0x7fd85e0905c0>

```



```
In [25]: from sklearn.base import BaseEstimator, RegressorMixin
from sklearn.utils.validation import check_X_y, check_array, check_is_fitted
from sklearn.linear_model import LinearRegression

class RegressionEnsembler(BaseEstimator, RegressorMixin):

    def __init__(self, estimators, learn_weights=True):
        self.estimators = estimators
        self.learn_weights = learn_weights

    def fit(self, X, y):

        X, y = check_X_y(X, y)

        self.X_ = X
        self.y_ = y

        for estimator in self.estimators:
            estimator.fit(X, y)

        if self.learn_weights:
            predictions = []
            for estimator in self.estimators:
                predictions.append(estimator.predict(X))
```



```

        Y_ = np.column_stack(predictions)

        self.ensembler = LinearRegression(fit_intercept=False)
        self.ensembler.fit(Y_, y)

    return self

def predict(self, X):
    check_is_fitted(self, ['X_', 'y_'])

    X = check_array(X)

    predictions = []
    for estimator in self.estimators:
        predictions.append(estimator.predict(X))
    Y_ = np.column_stack(predictions)

    if self.learn_weights:
        return self.ensembler.predict(Y_)
    else:
        return Y_.mean(axis=1)

```

```
In [26]: ensemble_model = RegressionEnsembler([model1, model3])
```

```

ensemble_model.fit(X_train, y_train)
ensemble_model.predict(X_train);

```

```
In [27]: -cross_val_score(RegressionEnsembler([model1, model2], learn_weights=False),
                        X_train, y_train,
                        cv=20, scoring="neg_mean_squared_error").mean()
```

```
Out[27]: 2.1874670566502465
```

```
In [28]: -cross_val_score(RegressionEnsembler([model1, model2], learn_weights=True),
                        X_train, y_train,
                        cv=20, scoring="neg_mean_squared_error").mean()
```

```
Out[28]: 2.1574493840185616
```

The Ensemble RMSE is actually greater than both individual models. In comparison to the 10 nearest neighbors model using the review text, the difference in RMSE's is negligible. The Ensemble RMSE is .03 points greater without learning the weights and is only .01 points better when learning the weights. However, the 80 nearest neighbors model using the foodie level number yielded an RMSE of 1.41 which is a bit less than the Ensemble model.

## 8 Submission Instructions

Once you are finished, follow these steps:

1. Restart the kernel and re-run this notebook from beginning to end by going to `Kernel > Restart Kernel and Run All Cells`. (If you are close to your API quota limit, do not re-run the code for Questions 1-3.)
2. If this process stops halfway through, that means there was an error. Correct the error and repeat Step 1 until the notebook runs from beginning to end.
3. Double check that there is a number next to each code cell and that these numbers are in order.

Then, submit your exam as follows:

1. Go to `File > Export Notebook As > PDF`.
2. Double check that the entire notebook, from beginning to end, is in this PDF file. (If the notebook is cut off, try first exporting the notebook to HTML and printing to PDF.)
3. Upload the PDF [to PolyLearn](#).